

Bertinelli_Gabriele_rlab01

April 9, 2024

1 RLab01 - Gabriele Bertinelli (2103359)

```
[ ]: library(tidyverse)
library(ggplot2)
library(tibble)
library(dplyr)
library(geosphere)
```

1)

```
[ ]: # read_csv is from tidyverse, faster than read.csv

d1 <- read_csv('./Data_CitiBike/JC-201902-citibike-tripdata.csv')
d2 <- read_csv('./Data_CitiBike/JC-201903-citibike-tripdata.csv')
d3 <- read_csv('./Data_CitiBike/JC-201904-citibike-tripdata.csv')
d4 <- read_csv('./Data_CitiBike/JC-201905-citibike-tripdata.csv')
d5 <- read_csv('./Data_CitiBike/JC-201906-citibike-tripdata.csv')
```

2)

```
[3]: # bind_rows() is useful for combining data frames wherein the columns are all
      ↪ identical
df <- bind_rows(d1, d2, d3, d4, d5)
```

```
[4]: str(df)
```

```
spc_tbl_ [150,792 × 15] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
 $ tripduration      : num [1:150792] 142 223 106 370 315 ...
 $ starttime         : POSIXct[1:150792], format: "2019-02-01 15:35:02"
"2019-02-01 17:00:46" ...
 $ stoptime          : POSIXct[1:150792], format: "2019-02-01 15:37:24"
"2019-02-01 17:04:30" ...
 $ start station id   : num [1:150792] 3183 3183 3183 3183 3183 ...
 $ start station name : chr [1:150792] "Exchange Place" "Exchange Place"
"Exchange Place" "Exchange Place" ...
 $ start station latitude : num [1:150792] 40.7 40.7 40.7 40.7 40.7 ...
 $ start station longitude: num [1:150792] -74 -74 -74 -74 -74 ...
 $ end station id      : num [1:150792] 3639 3681 3184 3211 3273 ...
```

```

$ end station name      : chr [1:150792] "Harborside" "Grand St" "Paulus Hook"
"Newark Ave" ...
$ end station latitude  : num [1:150792] 40.7 40.7 40.7 40.7 40.7 ...
$ end station longitude : num [1:150792] -74 -74 -74 -74 -74 ...
$ bikeid                : num [1:150792] 29677 26234 29588 29250 29586 ...
$ usertype              : chr [1:150792] "Subscriber" "Subscriber"
"Subscriber" "Subscriber" ...
$ birth year            : num [1:150792] 1963 1992 1960 1976 1980 ...
$ gender                : num [1:150792] 1 2 1 1 1 1 1 1 1 1 ...
- attr(*, "spec")=
.. cols(
..   tripduration = col_double(),
..   starttime = col_datetime(format = ""),
..   stoptime = col_datetime(format = ""),
..   `start station id` = col_double(),
..   `start station name` = col_character(),
..   `start station latitude` = col_double(),
..   `start station longitude` = col_double(),
..   `end station id` = col_double(),
..   `end station name` = col_character(),
..   `end station latitude` = col_double(),
..   `end station longitude` = col_double(),
..   bikeid = col_double(),
..   usertype = col_character(),
..   `birth year` = col_double(),
..   gender = col_double()
.. )
- attr(*, "problems")=<externalptr>

```

3)

```
[5]: df <- drop_na(df) # drop rows with NA values
```

```
[6]: str(df)
```

```
# no rows removed due to NA values
```

```

tibble [150,792 × 15] (S3: tbl_df/tbl/data.frame)
 $ tripduration      : num [1:150792] 142 223 106 370 315 ...
 $ starttime         : POSIXct[1:150792], format: "2019-02-01 15:35:02"
"2019-02-01 17:00:46" ...
 $ stoptime          : POSIXct[1:150792], format: "2019-02-01 15:37:24"
"2019-02-01 17:04:30" ...
 $ start station id   : num [1:150792] 3183 3183 3183 3183 3183 ...
 $ start station name : chr [1:150792] "Exchange Place" "Exchange Place"
"Exchange Place" "Exchange Place" ...
 $ start station latitude : num [1:150792] 40.7 40.7 40.7 40.7 40.7 ...
 $ start station longitude: num [1:150792] -74 -74 -74 -74 -74 ...

```

```

$ end station id      : num [1:150792] 3639 3681 3184 3211 3273 ...
$ end station name    : chr [1:150792] "Harborside" "Grand St" "Paulus Hook"
"Newark Ave" ...
$ end station latitude : num [1:150792] 40.7 40.7 40.7 40.7 40.7 ...
$ end station longitude : num [1:150792] -74 -74 -74 -74 -74 ...
$ bikeid              : num [1:150792] 29677 26234 29588 29250 29586 ...
$ usertype            : chr [1:150792] "Subscriber" "Subscriber"
"Subscriber" "Subscriber" ...
$ birth year          : num [1:150792] 1963 1992 1960 1976 1980 ...
$ gender              : num [1:150792] 1 2 1 1 1 1 1 1 1 1 ...

```

4) 4.1)

```

[7]: avg_trip_duration <- mean(df$tripduration)

median_trip_duration <- median(df$tripduration)

sprintf("Average trip duration: %.2f sec", avg_trip_duration)
sprintf("Median trip duration: %.2f sec", median_trip_duration)

```

'Average trip duration: 768.64 sec'

'Median trip duration: 341.00 sec'

4.2)

```

[8]: min_trip_duration <- min(df$tripduration)
max_trip_duration <- max(df$tripduration)

sprintf("Min trip duration: %.2f sec", min_trip_duration)
sprintf("Max trip duration: %.2f sec", max_trip_duration)

sprintf('max in hours: %.2f h', max_trip_duration/3600)

```

'Min trip duration: 61.00 sec'

'Max trip duration: 1729020.00 sec'

'max in hours: 480.28 h'

There's a trip that's really really long ahah. Almost 20 days. Maybe some long excursion... (or, more likely, wrong data)

4.3)

```

[9]: df_clean <- df %>% filter(tripduration <= 3*3600)

```

```

[10]: avg_trip_duration_clean <- mean(df_clean$tripduration)

median_trip_duration_clean <- median(df_clean$tripduration)

```

```

sprintf("Average trip duration: %.2f sec", avg_trip_duration_clean)
sprintf("Median trip duration: %.2f sec", median_trip_duration_clean)

min_trip_duration_clean <- min(df_clean$tripduration)
max_trip_duration_clean <- max(df_clean$tripduration)

sprintf("Min trip duration: %.2f sec", min_trip_duration_clean)
sprintf("Max trip duration: %.2f sec", max_trip_duration_clean)

sprintf('max in hours: %.2f h', max_trip_duration_clean/3600)

```

'Average trip duration: 553.38 sec'

'Median trip duration: 340.00 sec'

'Min trip duration: 61.00 sec'

'Max trip duration: 10800.00 sec'

'max in hours: 3.00 h'


4.4)

```

[11]: options(repr.plot.width = 15, repr.plot.height = 7)
time_duration_hist <- ggplot(data=df_clean, aes(x=tripduration)) +
  geom_histogram(binwidth=150, fill='mediumseagreen',
  ↪color='ivory', alpha=0.7) +
  labs(title='Trip Duration Distribution', x='Trip_
  ↪Duration (sec)', y='Frequency') +
  # xlim(c(0, 3600)) +
  theme_minimal()

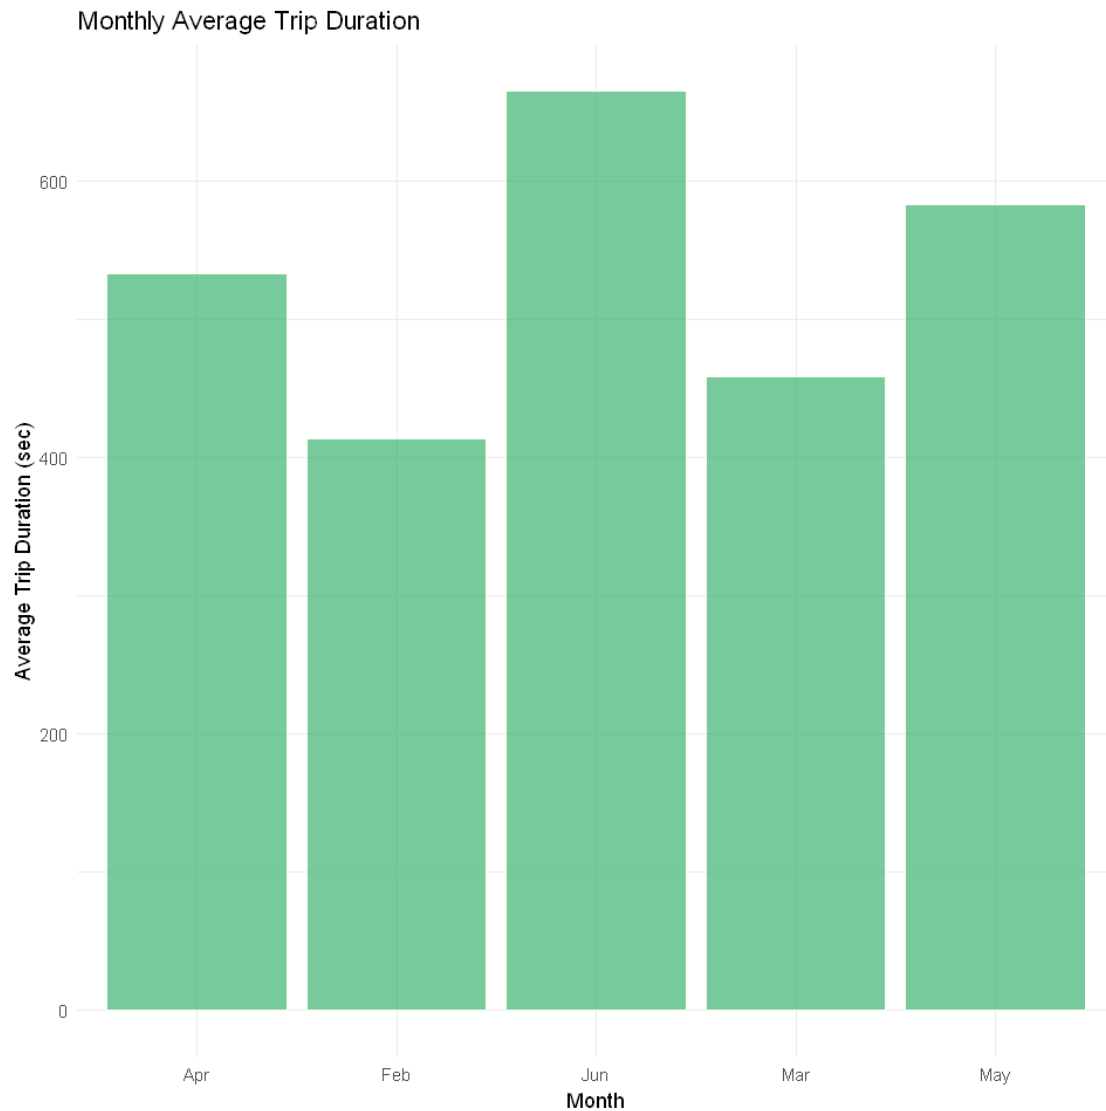
time_duration_hist

```



5)

```
[12]: monthly_avg_trip_duration <- df_clean %>%  
      mutate(month = lubridate::month(starttime)) %>% # mutate add a new col.␣  
      ↪ lubridate::month() returns the month number  
      group_by(month) %>%  
      summarize(avg_trip_duration = mean(tripduration)) # summarize() is used to␣  
      ↪ aggregate data by group  
  
monthly_avg_trip_duration <- monthly_avg_trip_duration %>% mutate(month_w =␣  
      ↪ c('Feb', 'Mar', 'Apr', 'May', 'Jun'))  
  
options(repr.plot.width = 8, repr.plot.height = 8)  
  
monthly_plot <- ggplot(data= monthly_avg_trip_duration, aes(x=month_w,␣  
      ↪ y=avg_trip_duration)) +  
      geom_bar(stat='identity', fill='mediumseagreen',␣  
      ↪ color='ivory', alpha=0.7) +  
      labs(title='Monthly Average Trip Duration', x='Month',␣  
      ↪ y='Average Trip Duration (sec)') +  
      theme_minimal()  
  
monthly_plot
```



6) 6.1)


```
[13]: riders_per_day <- df_clean %>%
mutate(date = as.Date(starttime)) %>%
group_by(date) %>%
summarize(riders = n())

options(repr.plot.width = 15, repr.plot.height = 7)

riders_per_day_plot <- ggplot(data=riders_per_day, aes(x=date, y=riders)) +
  geom_line(color='black', lwd=1) +
  labs(title='Ridership per Day', x='Date', y='Number of Riders') +
```


```
theme_minimal()
```

riders_per_day_plot



```
[14]: options(repr.plot.width = 15, repr.plot.height = 7)
riders_per_day_barplot <- ggplot(data=riders_per_day, aes(x=date, y=riders)) +
  geom_bar(stat='identity', fill='mediumseagreen',
  ↪color='ivory', alpha=0.7) +
  labs(title='Ridership per Day', x='Date', y='Number of
  ↪Riders') +
  theme_minimal()
```

riders_per_day_barplot



6.2)

```
[15]: # Extract the day of the week from the starttime column
df_clean$weekday <- lubridate::wday(df_clean$starttime, label = TRUE,
  ↪ locale="EN-us")


# Create a new column to differentiate between weekdays and weekends
df_clean$day_type <- ifelse(df_clean$weekday %in% c("Sat", "Sun"), "Weekend",
  ↪ "Weekday")

# Extract the hour from the starttime column
df_clean$hour <- hour(df_clean$starttime)

# Plot the hourly distribution on weekdays and weekends
options(repr.plot.width = 17, repr.plot.height = 7)

hourly_distribution <- ggplot(data = df_clean, aes(x = hour, fill = day_type)) +
  geom_histogram(binwidth = 1, position = "identity", alpha = 0.7,
  ↪ color='ivory') +
  labs(title = "Hourly Distribution on Weekdays and Weekends", x = "Hour", y =
  ↪ "Frequency") +
  scale_fill_manual(values = c("Weekday" = "mediumseagreen", "Weekend" =
  ↪ "navy")) +
  theme_minimal()

hourly_distribution
```



6.3)


```
[16]: # Filter the data for weekdays
df_weekdays <- df_clean %>% filter(day_type == "Weekday")

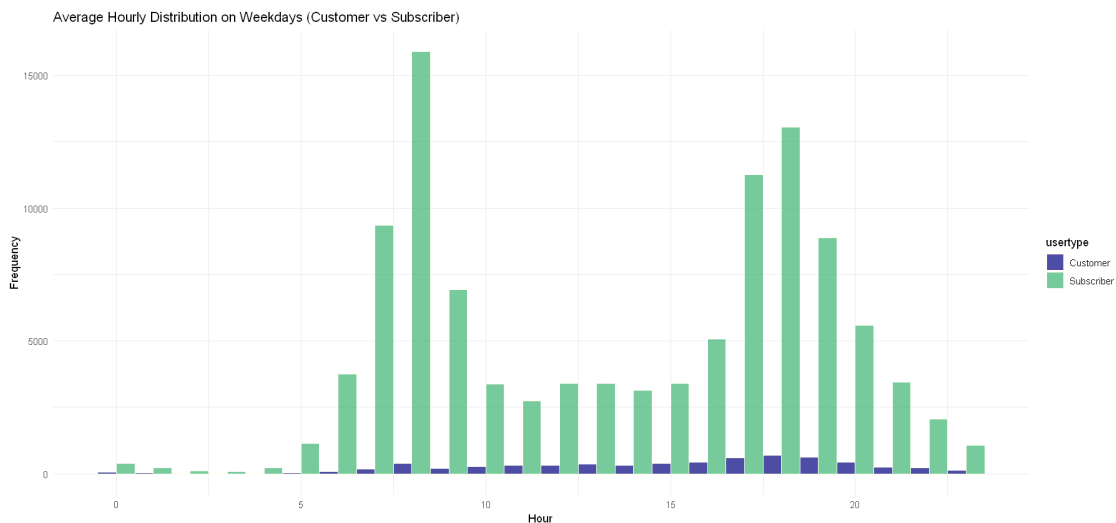
# Group the data by hour and user type
hourly_distribution <- df_weekdays %>%
  group_by(hour, usertype) %>%
  # summarize(avg_count = mean(tripduration)) %>%
  ungroup()

# Plot the average hourly distribution on weekdays, separating customer and
  ↳ subscriber users

options(repr.plot.width = 15, repr.plot.height = 7)

hourly_plot <- ggplot(data = hourly_distribution, aes(x = hour, fill =
  ↳ usertype)) +
  geom_histogram(binwidth = 1, position = "dodge", alpha = 0.7, color='ivory') +
  labs(title = "Average Hourly Distribution on Weekdays (Customer vs
  ↳ Subscriber)",
    x = "Hour", y = "Frequency") +
  scale_fill_manual(values = c("Customer" = "navy", "Subscriber" =
  ↳ "mediumseagreen")) +
  theme_minimal()

hourly_plot
```



7) 7.1)

```
[17]: df_vel <- df_clean %>% filter(tripduration > 3600)


df_vel$dist <- distHaversine(cbind(df_vel$start station longitude`,
  ↪df_vel$start station latitude`),
                                cbind(df_vel$end station longitude`, df_vel$end
  ↪station latitude`))

df_vel$speed <- (df_vel$dist / df_vel$tripduration)*3.6

options(repr.plot.width = 15, repr.plot.height = 7)

speed_hist <- ggplot(data=df_vel, aes(x=speed)) +
  geom_histogram(binwidth=0.1, fill='mediumseagreen',
  ↪color='ivory', alpha=0.7) +
  labs(title='Speed Distribution', x='Speed (m/s)',
  ↪y='Frequency') +
  theme_minimal()

speed_hist
```



7.2)

```
[18]: # Filter the data for the specified distance groups
df_group1 <- df_vel %>% filter(dist < 500)
df_group2 <- df_vel %>% filter(dist >= 500 & dist < 1000)
df_group3 <- df_vel %>% filter(dist >= 1000 & dist < 2000)
df_group4 <- df_vel %>% filter(dist >= 2000 & dist < 3000)
df_group5 <- df_vel %>% filter(dist >= 3000)
```

```

# Calculate the average speed for each distance group
avg_speed_group1 <- mean(df_group1$speed)
avg_speed_group2 <- mean(df_group2$speed)
avg_speed_group3 <- mean(df_group3$speed)
avg_speed_group4 <- mean(df_group4$speed)
avg_speed_group5 <- mean(df_group5$speed)

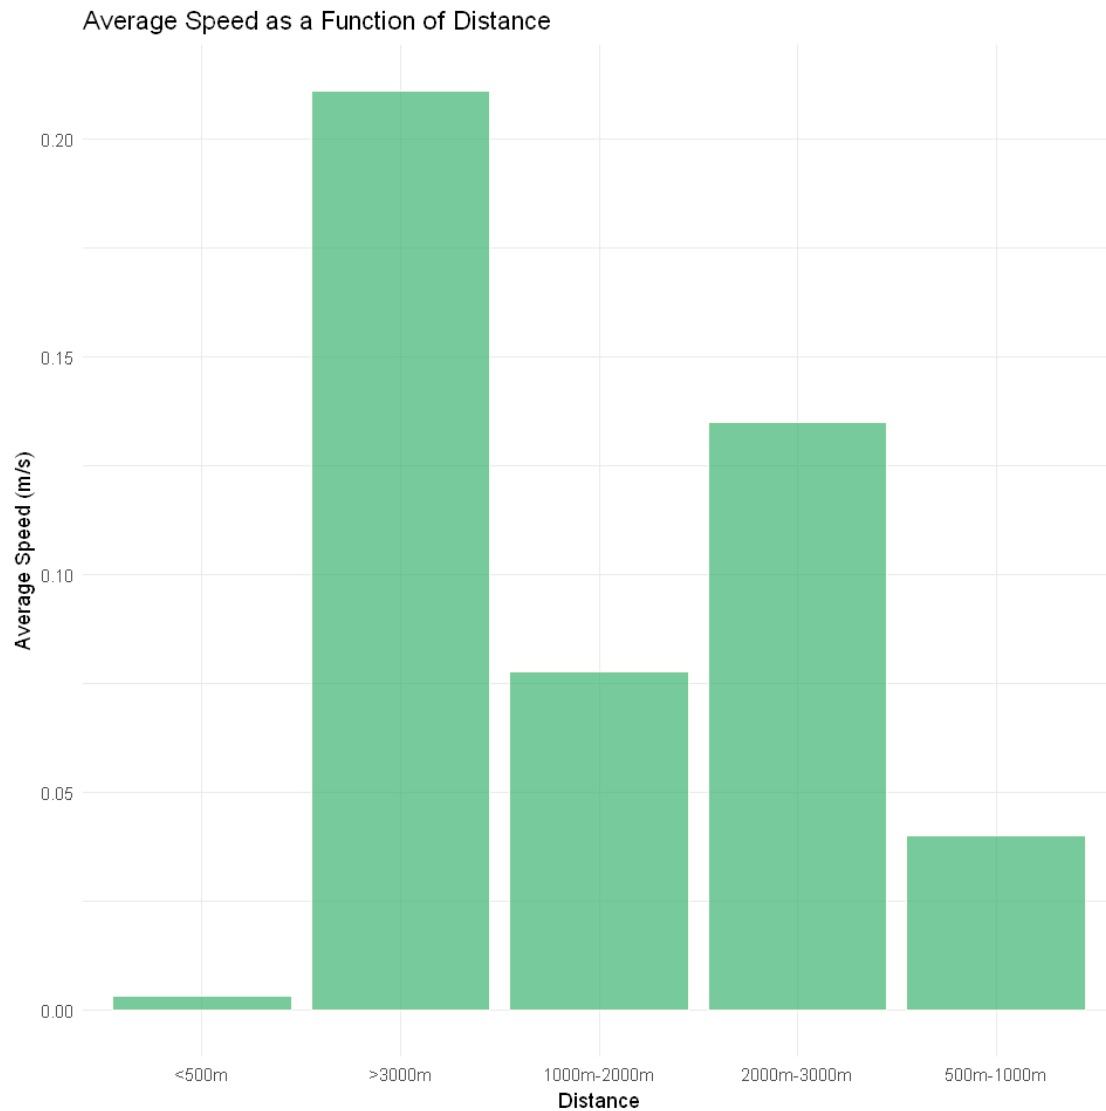
# Create a data frame for plotting
avg_speed_dist <- data.frame(group_dist = c("<500m", "500m-1000m",
↪ "1000m-2000m", "2000m-3000m", ">3000m"),
                             avg_speed = c(avg_speed_group1,
↪ avg_speed_group2, avg_speed_group3,
                             avg_speed_group4,
↪ avg_speed_group5))

options(repr.plot.width = 8, repr.plot.height = 8)

# Plot the average speed as a function of distance
speed_distance_plot <- ggplot(data = avg_speed_dist, aes(x = group_dist, y =
↪ avg_speed/3.6)) +
  geom_bar(stat = "identity", fill = "mediumseagreen",
↪ color = "ivory", alpha = 0.7) +
  labs(title = "Average Speed as a Function of Distance",
x = "Distance", y = "Average Speed (m/s)") +
  theme_minimal()

speed_distance_plot

```



7.3)

```
[19]: # Filter the data for weekdays
df_weekday <- df_vel %>% filter(day_type == "Weekday")

# Filter the data for weekends
df_weekend <- df_vel %>% filter(day_type == "Weekend")

# Now you can perform the same operations on df_weekday and df_weekend as before
# For example, for weekdays:

# Filter the data for the specified distance groups
df_group1d <- df_weekday %>% filter(dist < 500)
```

```

df_group2d <- df_weekday %>% filter(dist >= 500 & dist < 1000)
df_group3d <- df_weekday %>% filter(dist >= 1000 & dist < 2000)
df_group4d <- df_weekday %>% filter(dist >= 2000 & dist < 3000)
df_group5d <- df_weekday %>% filter(dist >= 3000)

df_group1w <- df_weekend %>% filter(dist < 500)
df_group2w <- df_weekend %>% filter(dist >= 500 & dist < 1000)
df_group3w <- df_weekend %>% filter(dist >= 1000 & dist < 2000)
df_group4w <- df_weekend %>% filter(dist >= 2000 & dist < 3000)
df_group5w <- df_weekend %>% filter(dist >= 3000)

# Calculate the average speed for each distance group
avg_speed_group1d <- mean(df_group1d$speed)
avg_speed_group2d <- mean(df_group2d$speed)
avg_speed_group3d <- mean(df_group3d$speed)
avg_speed_group4d <- mean(df_group4d$speed)
avg_speed_group5d <- mean(df_group5d$speed)

avg_speed_group1w <- mean(df_group1w$speed)
avg_speed_group2w <- mean(df_group2w$speed)
avg_speed_group3w <- mean(df_group3w$speed)
avg_speed_group4w <- mean(df_group4w$speed)
avg_speed_group5w <- mean(df_group5w$speed)

# Create a data frame for plotting
avg_speed_distd <- data.frame(group_dist = c("<500m", "500m-1000m", "1000m-2000m", "2000m-3000m", ">3000m"),
                              avg_speed = c(avg_speed_group1d, avg_speed_group2d, avg_speed_group3d, avg_speed_group4d, avg_speed_group5d),
                              day_type = 'Weekday')

avg_speed_distw <- data.frame(group_dist = c("<500m", "500m-1000m", "1000m-2000m", "2000m-3000m", ">3000m"),
                              avg_speed = c(avg_speed_group1w, avg_speed_group2w, avg_speed_group3w, avg_speed_group4w, avg_speed_group5w),
                              day_type = 'Weekend')

df_avg_speed_dist <- rbind(avg_speed_distd, avg_speed_distw)

options(repr.plot.width = 8, repr.plot.height = 8)

# Plot the average speed as a function of distance

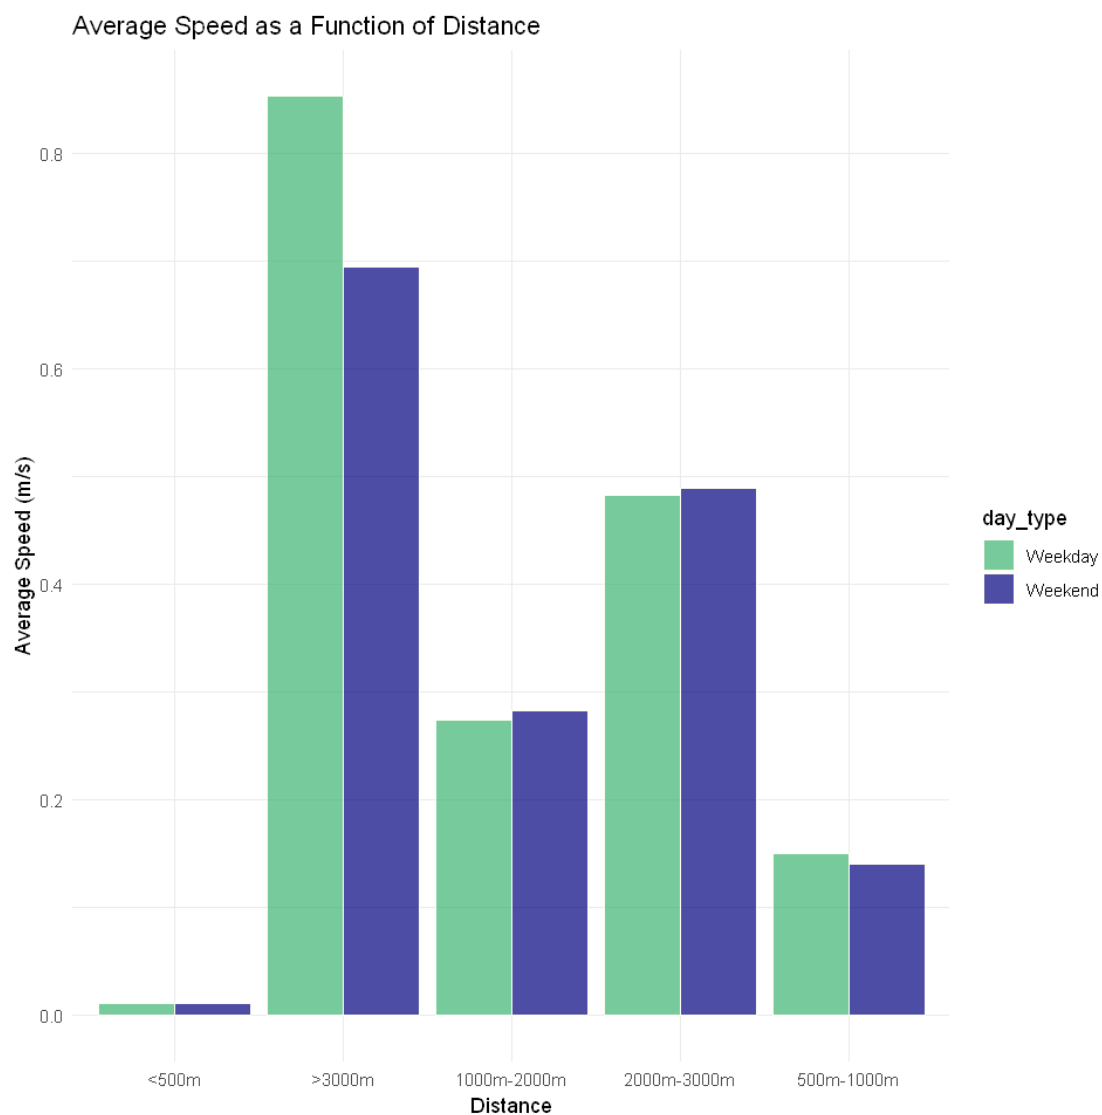
```

```

speed_distance_plotwd <- ggplot(data = df_avg_speed_dist, aes(x = group_dist, y =
  ↪ avg_speed, fill=day_type)) +
  geom_bar(position='dodge', stat = "identity", alpha =
  ↪ 0.7, color='ivory') +
  labs(title = "Average Speed as a Function of
  ↪ Distance",
        x = "Distance", y = "Average Speed (m/s)") +
  scale_fill_manual(values = c("Weekday" =
  ↪ "mediumseagreen", "Weekend" = "navy")) +
  theme_minimal()

```

speed_distance_plotwd



8) 8.1)


```
[20]: start_station <- df_clean %>% group_by(`start station name`) %>%  
      ↪ summarize(count = n()) %>% arrange(desc(count))  
  
end_station <- df_clean %>% group_by(`end station name`) %>% summarize(count =  
      ↪ n()) %>% arrange(count)  
  
sprintf('Most common start station --> %s', start_station$`start station`  
      ↪ name`[1])  
sprintf('Least popular end station --> %s', end_station$`end station name`[1])
```

'Most common start station -> Grove St PATH'

'Least popular end station -> 1 Ave & E 16 St'

8.2)

```
[21]: start_station <- df_clean %>% group_by(`start station name`) #>%  
      ↪ summarize(count = n())  
  
# start_station  
  
options(repr.plot.width = 15, repr.plot.height = 7)  
  
# start_distr <- ggplot(data=start_station, aes(x=`start station name`)) +  
#       geom_histogram(stat='count', binwidth = 1,  
#       ↪ fill='mediumseagreen', color='ivory', alpha=0.7) +  
#       labs(title='Start Station Distribution', x='Start Station',  
#       ↪ y='Frequency')  
#       theme_minimal()+  
#       theme(axis.text.x = element_text(angle = 45, hjust = 1))  
  
start_distr <- ggplot(data=start_station, aes(x=`start station name`)) +  
      geom_bar(fill='mediumseagreen', color='ivory', alpha=0.7) +  
      labs(title='Start Stations Distribution', x='Start Station',  
      ↪ y='Frequency') +  
      theme_minimal() +  
      theme(axis.text.x = element_text(angle = 45, hjust = 1))  
  
start_distr
```



8.3)

```
[22]: common_route <- df_clean %>% group_by(`start station name`, `end station name`)
      %>% summarize(count = n()) %>% arrange(desc(count))

least_route <- df_clean %>% group_by(`start station name`, `end station name`)
      %>% summarize(count = n()) %>% arrange(count)

cat('\n3 most common routes:\n')
for (i in 1:3) {
  print(sprintf('%s --> %s', common_route$start station name`[i],
    common_route$end station name`[i]))
}

cat('\n3 least popular routes:\n')
for (i in 1:3) {
  print(sprintf('%s --> %s', least_route$start station name`[i],
    least_route$end station name`[i]))
}
```

`summarise()` has grouped output by 'start station name'. You can override using the `.groups` argument.

`summarise()` has grouped output by 'start station name'. You can override using the `.groups` argument.

3 most common routes:

```
[1] "Hamilton Park --> Grove St PATH"
```


[1] "Grove St PATH --> Hamilton Park"
[1] "Brunswick & 6th --> Grove St PATH"

3 least popular routes:

[1] "5 Corners Library --> Dixon Mills"
[1] "5 Corners Library --> Grand St"
[1] "Astor Place --> Brunswick & 6th"